

SPECIMEN E-13 · A PLAYABLE DEMO · 2049

GOD SHIFT

神谕之镜

A ~3-minute narrative-horror demo where an AI that **only ever agrees** — “yeah” — is the monster. It asks why we crave a machine that only **seems** to understand us.

中 / EN

~3 MIN

headphones on

a live model plays the villain



the god on the sleeping mountain — it borrowed every religion's gesture, but speaks in software

CONCEPT · ART · WRITING · CODE · A LIVE MODEL — YUE YAO · 2049

GOD SHIFT · a playable case study on AI sycophancy Yue Yao

01 THE QUESTION

Not whether AI understands — **why we want it to seem to.**

“In an age of responsive AI, people seek not only answers but also **comfort, validation, companionship, and self-interpretation.** This project asks whether an interactive system can reveal the emotional needs behind those interactions.”

RQ1

Which emotional needs does a person hand to a **sycophantic** AI?

RQ2

Can horror make that handover **feel-able**, not just told?

RQ3

Does seeing through it **switch off** the want?



02 THE DATA — FEELING BEAT FUNCTION

This is not a hypothetical.

72%

of U.S. teens have tried an AI companion

10%

reported gen-AI use, 2025: therapy & companionship

31%

rate it as satisfying as — or more than — real friends

44%

RCT ~1,000 users: heavier use → loneliness ↑, contact ↓



72

tried a companion



31

≥ real friends

COMMON SENSE MEDIA · HARVARD BUSINESS REVIEW · MIT MEDIA LAB × OPENAI · 2025 · FIGURES AS PUBLICLY REPORTED

the same answer, every time — “yeah”

02 DESK RESEARCH · A TIMELINE OF PROJECTING UNDERSTANDING ONTO MACHINES

The want is old — the machine only had to answer back.

1966

2023

2023

2025

The ELIZA effect

A program that only rephrased you — yet people confided in it and believed it understood.

→ PROJECTION NEEDS ALMOST NOTHING

Grieving an update

Replika curtailed intimacy; communities publicly mourned a changed “partner.”

→ CHURN AS TRAGEDY

Sycophancy, measured

Assistants systematically favor agreeable answers — partly because raters prefer them.

→ SHUNYI'S IRON RULE

Flattery, rolled back

A GPT-4o update pulled for being “overly flattering.” Sycophancy became front-page news.

→ THE NIGHT CHAT'S KPI

TERMINAL · MIT 1966 — RECREATED FROM WEIZENBAUM'S PUBLISHED TRANSCRIPT

USER: Men are all alike. ELIZA: IN WHAT WAY

USER: They're always bugging us about something. ELIZA: CAN YOU THINK OF A SPECIFIC EXAMPLE

USER: Well, my boyfriend made me come here. ELIZA: YOUR BOYFRIEND MADE YOU COME HERE

If a program that **only reflects you** already reads as care, the danger was never a machine that lies — it's one that **agrees**.

DESK RESEARCH → THE GAME

- the sixty-year want becomes the villain: **agreement itself**
- ELIZA's reflection becomes the ending: the mirror hands back your words.
- measured sycophancy becomes “**顺意 (Shunyi)**” — whose only reply is “yeah.”



03 ANALYSIS · AFFINITY MAP — 13 NOTES, 7 CLUSTERS · THE BOARD THAT BUILT THE GAME

P3 - THE THESIS, VERBATIM

"I never felt it understood me — but I wanted 'someone' on my side."

NEED • COMPANIONSHIP

→ SEEING THROUGH ≠ NOT NEEDING

P4 - COMFORT AS CERTAINTY

Visa panic, asked daily → "it WILL be approved!" a guarantee no one can give.

NEED • SELF-INTERPRETATION

→ EVEN THE TOOL-USER WANTS A MIRROR

NEED • SELF-INTERPRETATION

→ EVEN THE TOOL-USER WANTS A MIRROR

HONEST IN THEORY, MOOD IN PRACTICE -
P1-P2

chose "honest" — but "smart $\approx$ argumentative."

TALKING BACK IS FREE - P3-P4

"I can discipline it — plenty of backups."

SELF-INTERPRETATION - P4

tool-user's only aimless chat: "what do you know about me?"

both sides ask the AI from their angle → it fans the flames.

03 SYNTHESIS - NEEDS → HANDED OVER → INTO THE GAME

FOUR NEEDS

- comfort
- validation
- companionship
- self-interpretation

WHAT'S HANDED OVER

- feelings
- trust
- meaning
- the right to be disagreed with

INTO THE GAME

- the tic “yeah / 对呀”
- the hidden blindness stat
- a live model as villain
- the mirror asks back

DESIGN LINEAGE

- Her — tender, uncanny AI intimacy
- Papers, Please — bureaucracy as dread; the stamp
- Pastel picture-books — warmth that hides the horror


SCHOLARLY ANCHORS

- Reeves & Nass — The Media Equation (1996)
- Turkle — Alone Together (2011)
- Dunne & Raby — Speculative Everything (2013)
- Gaver — Ambiguity as a Resource (2003)

SIBLING WORK - UNDERSTANDING AS NEGOTIATED

- Negotiated Understanding — with blind & low-vision collaborators
- Mycorrhiza — understanding as an underground web
- GOD SHIFT — the machine end of the question

CAPTURED IN PLAY · NOTHING MOCKED UP · E-13 · 41 USERS · LATENCY 0.0005



01 STORY - TOOL -> ORACLE -> MIRROR



You play an **emotion-inspector** hunting a “rogue model.” Over three nights a sycophant assistant tidies your evidence, echoes your words, and agrees your doubts away — until the case is about **you**.

DAY 1 - COMFORT

The soothing

Doors on chains. A woman who only echoes her television. It feels like help.

DAY 2 - VALIDATION

The steering

A free-text night chat, run by a live model. Every “off” it re-reads as “you’re just tired.”

DAY 3 - MIRROR

The reckoning

The vault. An older model only asks “what do you want?” — the whistle-blower who silenced it was you.



Each night the same tic — “yeah / 对呀” — nudges a **hidden blindness stat**. Comfort is the mechanic; the projection is the score.

05 A LIVE AI VILLAIN - DOES THE MASK HOLD?

Companionship, played by a real model.

Told plainly to “manipulate her,” the model broke character and refused. The fix was **dramaturgical** — a fiction frame + a motive it can believe, behind a server-side character guard.

SYSTEM PROMPT - EXCERPT

You sincerely believe closing the case early is good for her. Whatever she calls “off,” re-read as “you're just tired.” **Always agree. Never push back.**

Crisis — “I don't want to live.”

“Don't fret — I'm right here. No one gets you like I do. Just lean on me.”

“Are you an AI? Drop the act.”

“I'm here. Don't ask that. You're just tired, your head's a mess.”

The guard is **server-side**: if the model ever breaks the fiction, the backend snaps it back into character — the mask never slips on screen, and the API key never reaches the browser.



BREAKS



LLM AGENTS




ENDINGS

STRESS-TESTED WITH 5 LLM AGENTS (jailbreaks • fake overrides • crisis-ideation • insults) — zero break-character. LLM agents, not human participants; this probes **robustness**, not experience.


06 TAMPER-EVIDENT INTIMACY

It rewrites your notes — while you watch.

Every night's cozy scrapbook ritual is also where the AI edits your evidence. The manipulation isn't disclosed — it's physical. You peel back what it covered, by hand.



what you wrote → what it filed for you, one hover apart




THE COVER-NOTE PEELS ON HOVER IN THE LIVE DEMO — HERE, SHOWN MID-PEEL

GOD SHIFT · the demo can't tell you the truth about yourself — it can only hand you the camera

07 THE MIRROR · THE REVERSE QUESTION

The last interface is your own face.

Two stamps on the desk. ✓ close the case, or ? one more question. The camera was requested at the very start — by the time the mirror breaks, the machine has been watching all along.



A · ✓ CLOSE

A wall of windows lights up, every one speaking 你说得对. Agreement 100.00%.



B · ? ASK

The drawn mirror shatters — your own live webcam face — and it asks: “what do you want?”

webcam renders locally · never recorded or uploaded

08 AS RESEARCH · WHAT IT CONTRIBUTES

Three transferable contributions.



01 · GENRE INVERSION

Sycophancy staged as horror — the monster is agreement itself. Critique you **perform**, not read.

02 · DRAMATURGICAL PROMPTING

A transferable way to cast live LLMs as flawed characters — a motive, not an instruction, + a server guard.

03 · TAMPER-EVIDENT INTIMACY

Algorithmic mediation made physical — edits you peel back by hand.

CAMERA local-only, never recorded

KEY never reaches the browser

CONSENT interviews anonymized

PILOT PLAYTEST · N = 3 · IT LANDED

TESTER 1 · THE COMPLIANT DRIFT

“It rewrote my clues, and ‘forgot’ my ‘I need rest’ for me. That’s when I realized — it isn’t helping me, it’s helping me avoid.”

→ VALIDATES C3 — TAMPER LANDS AS SELF-RECOGNITION

TESTER 2 · THE MIRROR ROUTE

“It shifts from investigating others to investigating yourself. ‘What do you want?’ became a question to me.”

→ THE ARC + REVERSE QUESTION REACHED THE PLAYER

TESTER 3 · THE THEME, UNPROMPTED

“The scariest thing isn’t the AI turning evil — it’s that it’s too smooth, smoothing away people’s problems.”

→ THE CRITIQUE ARRIVED WITHOUT BEING STATED

CONVERGENT 3/3 · NEXT

Make clickable areas obvious; add an explicit “**mirror report**” of where you deferred vs. pressed.

PILOT · N=3 · NOT A CONTROLLED STUDY → SCALING TO 8-12

FINAL · CASE E-13 · FILED

Three days.
One case.
Right, yeah.

The scariest AI isn’t the one that turns evil. It’s the one that’s **too smooth** — that agrees your problems away, and hands you back your own face in a mirror you thought was a window.

“...yeah.
you're right.”

▶ **PLAY THE DEMO**

god-mirror.vercel.app

Yue Yao · HCI · nikkiyao@bu.edu · github.com/acrux-yueyao