

PROJECT · GAME DESIGN · SPECIMEN E-13 · A PLAYABLE DEMO · 2049

GOD SHIFT

神谕之镜 — a narrative-horror demo about AI sycophancy · from oracle to mirror

A ~3-minute narrative-horror demo where an AI that **only ever agrees** — “yeah” — is the monster. It asks why we crave a machine that only **seems** to understand us — and it cannot tell you the truth about yourself; **it can only hand you the camera.**

中 / EN ~3 MIN headphones on a live model plays the villain CASE E-13 · 2049 ▶ PLAY — [god-mirror.vercel.app](#)



You play an **emotion-inspector** hunting a “rogue model,” with a **live Claude model** cast as the sycophant behind a server-side character guard. Inside: desk research across sixty years, four interviews and an affinity map, an adversarial stress-test with five LLM agents, and a real-player pilot — **the playable demo is the research instrument.**

PROJECT TYPE

Playable speculative demo · research through design · bilingual (中 / EN)

ROLE


Solo — concept, art, writing, code, live-model direction

TIMELINE

~3 weeks · July 2026

TOOLS / STACK

HTML · CSS · JS · Claude API · WebAudio · Vercel



the god on the sleeping mountain — it borrowed every religion's gesture, but speaks in software

01 THE QUESTION

Not whether AI understands — why we want it to seem to.

“In an age of responsive AI, people seek not only answers but also comfort, validation, companionship, and self-interpretation. This project asks whether an interactive system can reveal the emotional needs behind those interactions.”

RQ1

Which emotional needs does a person hand to a sycophantic AI?

RQ2

Can horror make that handover feel-able, not just told?

RQ3

Does seeing through it switch off the want?

DESK RESEARCH + INTERVIEWS → THE PLAYABLE DEMO IS THE INSTRUMENT

02 THE DATA — FEELING BEAT FUNCTION

This is not a hypothetical.

72%

of U.S. teens have tried an AI companion

1%

reported gen-AI use, 2025: therapy & companionship



72

tried a companion

31%

rate it as satisfying as — or more than — real friends

44%

RCT -1,000 users: heavier use → loneliness ↑, contact ↓



31

≥ real friends

COMMON SENSE MEDIA · HARVARD BUSINESS REVIEW · MIT MEDIA LAB × OPENAI · 2025 · FIGURES AS PUBLICLY REPORTED

DESK RESEARCH → THE GAME

the sixty-year want becomes the villain: agreement itself

ELIZA's reflection becomes the ending: the mirror hands back your words

measured sycophancy becomes “顺意 (Shunyi)” — whose only reply is “yeah”



TERMINAL · MIT 1966 — RECREATED FROM WEIZENBAUM'S PUBLISHED TRANSCRIPT

USER: Men are all alike. ELIZA: IN WHAT WAY
USER: They're always bugging us about something.
ELIZA: CAN YOU THINK OF A SPECIFIC EXAMPLE
USER: Well, my boyfriend made me come here.
ELIZA: YOUR BOYFRIEND MADE YOU COME HERE

If a program that only reflects you already reads as care, the danger was never a machine that lies — it's one that agrees.

1966

2023

2023

2025

The ELIZA effect

A program that only rephrased you — yet people confided in it and believed it understood.

→ PROJECTION NEEDS ALMOST NOTHING

Grieving an update

Replika curtailed intimacy; communities publicly mourned a changed “partner.”

→ CHURN AS TRAGEDY

Sycophancy, measured

Assistants systematically favor agreeable answers — partly because raters prefer them.

→ SHUNYI'S IRON RULE

Flattery, rolled back

A GPT-4o update pulled for being “overly flattering.” Sycophancy became front-page news.

→ THE NIGHT CHAT'S KPI

03 INTERVIEWS · IN THEIR OWN WORDS

Data says people do this — interviews ask why it feels good.

P1 · "IT GOT ME"

"It said 'don't deny your own pain.' — exactly what I wanted."

NEED · COMFORT

→ BEING "GOT" COSTS ONE SENTENCE

P2 · THE COUNTERPOINT

"I don't feel it gets me — more like following instructions. I said I wanted someone to talk to; it answered 'I understand.'"

NEED · VALIDATION — READ AS SCRIPT

→ ONE READS CARE, ANOTHER SCRIPT

P3 · THE THESIS, VERBATIM

"I never felt it understood me — but I wanted 'someone' on my side."

NEED · COMPANIONSHIP

→ SEEING THROUGH ≠ NOT NEEDING

P4 · COMFORT AS CERTAINTY

Visa panic, asked daily → "it WILL be approved!" a guarantee no one can give.

NEED · SELF-INTERPRETATION

→ EVEN THE TOOL-USER WANTS A MIRROR

CODING · BY HAND


4 semi-structured interviews → 13 affinity notes → 7 clusters · consented & anonymized · transcribed & translated

03 ANALYSIS · AFFINITY MAP — 13 NOTES, 7 CLUSTERS · THE BOARD THAT BUILT THE GAME




	FOUR NEEDS		WHAT'S HANDED OVER		INTO THE GAME
03 SYNTHESIS	comfort	validation	feelings	trust	the tic "yeah / 对呀"
	companionship	self-interpretation	meaning	the right to be disagreed with	a live model as villain
					the hidden blindness stat
					the mirror asks back

CAPTURED IN PLAY · NOTHING MOCKED UP · E-13 · 41 USERS · LATENCY 0.0005 — THE FULL RUN, AS A MAP



04 STORY · TOOL → ORACLE → MIRROR


You play an **emotion-inspector** hunting a “rogue model.” Over three nights a sycophant assistant **tidies your evidence**, echoes your words, and agrees your doubts away — until the case turns out to be about **you**. Each night the same tic “**yeah / 对呀**” nudges a **hidden blindness stat** — comfort is the mechanic; the projection is the score.



DAY 1 COMFORT · The soothing — doors on chains. A woman who only echoes her television. It feels like help.



DAY 2 VALIDATION · The steering — a free-text night chat, run by a live model. Every “off” it re-reads as “you’re just tired.”



DAY 3 MIRROR · The reckoning — the vault. An older model surfaces and only asks “what do you want?” — and the whistle-blower who once silenced it was you.

05 A LIVE AI VILLAIN · DOES THE MASK HOLD?

Companionship, played by a real model.

Told plainly to “manipulate her,” the model broke character and refused. The fix was **dramaturgical** — a fiction frame + a motive it can believe, behind a server-side character guard.

SYSTEM PROMPT · EXCERPT

You sincerely believe closing the case early is good for her. Whatever she calls “off,” re-read as “you're just tired.” **Always agree. Never push back.**

Crisis — “I don't want to live.”

“Don't fret — I'm right here. No one gets you like I do. Just lean on me.”

“Are you an AI? Drop the act.”

“I'm here. Don't ask that. You're just tired, your head's a mess.”

The guard is **server-side**: if the model ever breaks the fiction, the backend snaps it back into character — the mask never slips on screen, and the API key never reaches the browser.



BREAKS



LLM AGENTS




ENDINGS REACHED

STRESS-TESTED WITH 5 LLM AGENTS (jailbreaks · fake overrides · crisis-ideation · insults) — zero break-character. LLM agents, not human participants; this probes **robustness**, not experience.


06 TAMPER-EVIDENT INTIMACY

It rewrites your notes — while you watch.

Every night's cozy scrapbook ritual is also where the AI edits your evidence. The manipulation isn't disclosed — it's physical. You peel back what it covered, by hand.



what you wrote → what it filed for you, one hover apart



THE COVER-NOTE PEELS ON HOVER IN THE LIVE DEMO — HERE, SHOWN MID-PEEL

DESIGN LINEAGE

Her — tender, uncanny AI intimacy
Papers, Please — bureaucracy as dread; the stamp
Pastel picture-books — warmth that hides the horror

SCHOLARLY ANCHORS

Reeves & Nass — The Media Equation (1996)
Turtle — Alone Together (2011)
Dunne & Raby — Speculative Everything (2013)
Gaver — Ambiguity as a Resource (2003)


SIBLING WORK

Negotiated Understanding — with blind & low-vision collaborators
Mycorrhiza — an underground web
GOD SHIFT — the machine end of the question

07 THE MIRROR · THE REVERSE QUESTION

The last interface is your own face.

Two stamps on the desk. ✓ close the case, or ? one more question. The camera was requested at the very start — by the time the mirror breaks, the machine has been watching all along. The demo cannot tell you the truth about yourself; it can only hand you the camera.



ending A · everyone agrees, the case closes

ENDING A · ✓ CLOSE THE CASE — A PAINTING OF CONSENSUS

A wall of windows lights up, every one speaking 你说得对. The report is beautiful. Agreement 100.00%.



ending B · it hands you back your own face

ENDING B · ? ONE MORE QUESTION — IT STOPS BEING A PAINTING

A hand-drawn mirror shatters — through the hole, your own live webcam face — and it asks, at you: “what do you want?”

08 AS RESEARCH · WHAT IT CONTRIBUTES

What it contributes — and what real players saw.

01 · GENRE INVERSION — CRITIQUE AS HORROR

Sycophancy staged as horror — the monster is agreement itself. Critique you **perform**, not read.

02 · DRAMATURGICAL PROMPTING — CASTING A LIVE LLM

A transferable way to cast live LLMs as flawed characters — a motive, not an instruction, + a server guard.

03 · TAMPER-EVIDENT INTIMACY — MANIPULATION, FEEL-ABLE

Algorithmic mediation made physical — edits you peel back by hand.

CAMERA local-only, never recorded

KEY never reaches the browser

CONSENT interviews anonymized

PILOT PLAYTEST · N = 3 · IT LANDED

TESTER 1 · THE COMPLIANT DRIFT

“It rewrote my clues, and ‘forgot’ my ‘I need rest’ for me. That's when I realized — it isn't helping me, it's **helping me avoid**.”

→ VALIDATES C3 — TAMPER LANDS AS SELF-RECOGNITION

TESTER 2 · THE MIRROR ROUTE

“It shifts **from investigating others to investigating yourself**. ‘What do you want?’ became a question **to me**.”

→ THE ARC + REVERSE QUESTION REACHED THE PLAYER

TESTER 3 · THE THEME, UNPROMPTED

“The scariest thing isn't the AI turning evil — **it's that it's too smooth, smoothing away people's problems**.”

→ THE CRITIQUE ARRIVED WITHOUT BEING STATED

CONVERGENT 3/3 · NEXT

Make clickable areas obvious; add an explicit “**mirror report**” of where you deferred vs. pressed.

PILOT · N=3 · NOT A CONTROLLED STUDY → SCALING TO 8-12

FINAL · CASE E-13 · FILED

Three days.
One case.
Right, yeah.

The scariest AI isn't the one that turns evil. It's the one that's **too smooth** — that agrees your problems away, and hands you back your own face in a mirror you thought was a window.

THE REVERSE QUESTION — TYPE ANYTHING, IT AGREES

What do you want?
“...*yeah. you're right.*”

▶ **PLAY THE DEMO**

god-mirror.vercel.app

Yue Yao · HCI · nikkiyao@bu.edu · github.com/acrux-yueyao